

CMPE 259: Homework 1: Regular expressions, text normalization, and edit distance

The parts that you need to complete are marked as Exercises.

Part 0: Initialization & Setup

```
In [1]: # importing required libraries
import re
import nltk
from nltk.corpus import movie_reviews
import string
import pandas as pd
from nltk.corpus import stopwords

nltk.download('movie_reviews')
nltk.download('stopwords')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package movie_reviews to /root/nltk_data...
[nltk_data]   Unzipping corpora/movie_reviews.zip.
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
[nltk_data] Downloading package wordnet to /root/nltk_data...
```

```
Out[1]: True
```

Part 1: Regular Expressions

Extracting license plate numbers, IDs, emails and mailing addresses from a document

Document creation

```
In [2]: sentence = 'I am 20 years old. My previous license plate number was 4XUI302 and my new o
sentence
```

```
Out[2]: 'I am 20 years old. My previous license plate number was 4XUI302 and my new one is 3A-27
8. My ID is J987492 and my address is 123 Main street, San Jose, CA. Please email me at
myemail123+spam@google.cg or jane.doe@sjsu.edu'
```

Extracting license plate numbers

```
In [3]: # The format of license plate number is a digit then 2 or 3 letters (one of which can be

regex = re.compile(r'(\d{1}[A-Za-z-]{2,3}\d{3})')
license_plate_numbers = regex.findall(sentence)
license_plate_numbers
```

```
Out[3]: ['4XUI302', '3A-278']
```

Exercise 1-1: Extract the ID numbers from the document.

```
In [4]: # The format of the IDs is one character/letter and then 6 digits
        regex = re.compile(r'([A-Za-z]\d{6})')
        ids = regex.findall(sentence)
        ids
```

```
Out[4]: ['J987492']
```

Exercise 1-2: Extract the email IDs from the document

```
In [5]: regex = re.compile(r'([\d\w+.*@]+\w\d+[.]+[a-z]*)')
        emails = regex.findall(sentence)
        emails
```

```
Out[5]: ['myemail123+spam@google.cg', 'jane.doe@sjsu.edu']
```

Exercise 1-3: Extract the mailing address from the document

```
In [6]: # Starts with 3 digit number, then some letters, a comma, letters, a comma
        regex = re.compile(r'([0-9 ]{3}[A-Za-z ]*[,{1}[A-Za-z0-9 ]*[,{1}\w ]*)')
        mailing_address = regex.findall(sentence)
        mailing_address
```

```
Out[6]: ['123 Main street, San Jose, CA']
```

Exercise 1-4: Anonymize the license plate numbers by replacing them with the text "LP_NUM"

The re.sub function is described here: <https://docs.python.org/3/library/re.html>

```
In [9]: # Now replacing license plate numbers with the string "LP_NUM"
        sentence_modified = sentence
        for i in license_plate_numbers:
            sentence_modified = sentence_modified.replace(i, "LP_NUM")
        sentence_modified
```

```
Out[9]: 'I am 20 years old. My previous license plate number was LP_NUM and my new one is LP_NUM
        M. My ID is J987492 and my address is 123 Main street, San Jose, CA. Please email me at
        myemail123+spam@google.cg or jane.doe@sjsu.edu'
```

Exercise 1-5: Replace the ID numbers with the text "ID_NUM"

```
In [10]: for i in ids:
          sentence_modified = sentence_modified.replace(i, "ID_NUM")
          sentence_modified
```

```
Out[10]: 'I am 20 years old. My previous license plate number was LP_NUM and my new one is LP_NUM
        M. My ID is ID_NUM and my address is 123 Main street, San Jose, CA. Please email me at m
        yemail123+spam@google.cg or jane.doe@sjsu.edu'
```

Part 2: Text Processing

Count the number of words in the movie_reviews dataset (dataset uploaded in the beginning of this notebook under "Part 0: Initialization and Setup")

```
In [11]: # print number of words in the movie review dataset
        len(movie_reviews.words())
```

Out[11]: 1583820

Load the standard list of punctuation marks

```
In [12]: punctuations = string.punctuation
punctuations
```

Out[12]: '!"#\$%&\'()*+,-./:;<=>?@[\\]^_`{|}~'

Remove punctuation from movie reviews

```
In [13]: words_wo_puncts = [x for x in movie_reviews.words() if x not in punctuations]
len(words_wo_puncts)
```

Out[13]: 1338788

Count the number of unique words

```
In [14]: unique_words = set(words_wo_puncts)
len(unique_words)
```

Out[14]: 39737

Find the 20 most frequent words in the dataset

```
In [15]: # top 20 highest freq words
pd.Series(words_wo_puncts).value_counts()[:20]
```

Out[15]:

	count
--	-------

the	76529
-----	-------

a	38106
---	-------

and	35576
-----	-------

of	34123
----	-------

to	31937
----	-------

is	25195
----	-------

in	21822
----	-------

s	18513
---	-------

it	16107
----	-------

that	15924
------	-------

as	11378
----	-------

with	10792
------	-------

for	9961
-----	------

his	9587
-----	------

this	9578
------	------

film	9517
------	------

i	8889
---	------

he	8864
----	------

but	8634
-----	------

on	7385
----	------

dtype: int64

Load the standard list of stopwords

```
In [16]: # getting english stopwords
eng_stopwords = stopwords.words('english')
eng_stopwords
```

```
Out[16]: ['i',
'me',
'my',
'myself',
'we',
'our',
'ours',
'ourselves',
'you',
"you're",
"you've",
"you'll",
"you'd",
'your',
'yours',
'yourself',
'yourselves',
'he',
'him',
'his',
'himself',
'she',
"she's",
'her',
'hers',
'herself',
'it',
"it's",
'its',
'itself',
'they',
'them',
'their',
'theirs',
'themselves',
'what',
'which',
'who',
'whom',
'this',
'that',
"that'll",
'these',
'those',
'am',
'is',
'are',
'was',
'were',
'be',
'been',
'being',
'have',
'has',
'had',
'having',
```

'do',
'does',
'did',
'doing',
'a',
'an',
'the',
'and',
'but',
'if',
'or',
'because',
'as',
'until',
'while',
'of',
'at',
'by',
'for',
'with',
'about',
'against',
'between',
'into',
'through',
'during',
'before',
'after',
'above',
'below',
'to',
'from',
'up',
'down',
'in',
'out',
'on',
'off',
'over',
'under',
'again',
'further',
'then',
'once',
'here',
'there',
'when',
'where',
'why',
'how',
'all',
'any',
'both',
'each',
'few',
'more',
'most',
'other',
'some',
'such',
'no',
'nor',
'not',
'only',
'own',
'same',

```
'so',  
'than',  
'too',  
'very',  
's',  
't',  
'can',  
'will',  
'just',  
'don',  
'don't',  
'should',  
'should've',  
'now',  
'd',  
'll',  
'm',  
'o',  
're',  
've',  
'y',  
'ain',  
'aren',  
'aren't',  
'couldn',  
'couldn't',  
'didn',  
'didn't',  
'doesn',  
'doesn't',  
'hadn',  
'hadn't',  
'hasn',  
'hasn't',  
'haven',  
'haven't',  
'isn',  
'isn't',  
'ma',  
'mightn',  
'mightn't',  
'mustn',  
'mustn't',  
'needn',  
'needn't',  
'shan',  
'shan't',  
'shouldn',  
'shouldn't',  
'wasn',  
'wasn't',  
'weren',  
'weren't',  
'won',  
'won't',  
'wouldn',  
'wouldn't']
```

Count the number of stopwords

```
In [17]: len(eng_stopwords)
```

```
Out[17]: 179
```

Exercise 2-1: Remove the stopwords from the dataset (similarly to how we removed punctuation above)

```
In [18]: words_wo_puncts_stopwords = [i for i in words_wo_puncts if i not in eng_stopwords]
len(words_wo_puncts_stopwords)
```

Out[18]: 710578

Exercise 2-2: Find the number of unique words in the dataset now that the stop words have been removed

```
In [19]: # unique words without stopwords
unique_words = set(words_wo_puncts_stopwords)
len(unique_words)
```

Out[19]: 39586

Exercise 2-3: Find the top 20 highest frequency words now that we have removed the stopwords

```
In [21]: # top 20 highest freq words after removing stopwords
pd.Series(words_wo_puncts_stopwords).value_counts()[:20]
```

Out[21]:

	count
film	9517
one	5852
movie	5771
like	3690
even	2565
time	2411
good	2411
story	2169
would	2109
much	2049
character	2020
also	1967
get	1949
two	1911
well	1906
characters	1859
first	1836
--	1815
see	1749
way	1693

dtype: int64

Find the words that are used only once in the corpus (and print the first few).

```
In [22]: # 20 words that are used only once in corpus using hapaxes() function
         nltk.FreqDist(words_wo_puncts_stopwords).hapaxes()[:20]
```

```
Out[22]: ['loooooot',
          'schnazzy',
          'timex',
          'indiglo',
          'jessalyn',
          'gilsig',
          'ruber',
          'jaleel',
          'balki',
          'wavers',
          'statistics',
          'snapshot',
          'guesswork',
          'maryam',
          'daylights',
          'terraformed',
          'stagnated',
          'napolean',
          'millimeter',
          'enmeshed']
```

Exercise 2-4: Use the PorterStemmer to stem the words in the dataset.

Display the first few words.

```
In [30]: from nltk.stem import PorterStemmer

         ps = PorterStemmer()
         words_wo_puncts_stopwords_stemmed = [ps.stem(i) for i in words_wo_puncts_stopwords]

         for i in range(20):
             print(words_wo_puncts_stopwords[i], ": ", words_wo_puncts_stopwords_stemmed[i])

plot : plot
two : two
teen : teen
couples : coupl
go : go
church : church
party : parti
drink : drink
drive : drive
get : get
accident : accid
one : one
guys : guy
dies : die
girlfriend : girlfriend
continues : continu
see : see
life : life
nightmares : nightmar
deal : deal
```

Exercise 2-5: Use the WordNetLemmatizer to lemmatize the words in the dataset.

Display the first few words.


```
In [31]: from nltk import WordNetLemmatizer

lemmatizer = WordNetLemmatizer()

words_wo_puncts_stopwords_lemma = [lemmatizer.lemmatize(i) for i in words_wo_puncts_stopwords]

for i in range(20):
    print(words_wo_puncts_stopwords[i], ": ", words_wo_puncts_stopwords_lemma[i])

plot : plot
two : two
teen : teen
couples : couple
go : go
church : church
party : party
drink : drink
drive : drive
get : get
accident : accident
one : one
guys : guy
dies : dy
girlfriend : girlfriend
continues : continues
see : see
life : life
nightmares : nightmare
deal : deal
```

Exercise 2-6:

- How many unique words are there once stemming is applied? (show the code that performs the computation and outputs the result)
- How many unique words are there once lemmatization is applied? (show the code that performs the computation and outputs the result)

```
In [32]: print("Unique words after stemming: ", len(set(words_wo_puncts_stopwords_stemmed)))
print("Unique words after lemmatization: ", len(set(words_wo_puncts_stopwords_lemma)))

Unique words after stemming: 26101
Unique words after lemmatization: 35172
```

Part 3. Tokenization

Exercise 3-1: Use the Penn Tree Bank tokenizer to tokenize the sentence below

Print the tokens that the tokenizer produces.

```
In [33]: from nltk.tokenize import TreebankWordTokenizer

s = 'Please pay $100.55 to settle your bill. Send confirmation to confirm@gmail.com.'

tokenizer = TreebankWordTokenizer()
s_tokens = tokenizer.tokenize(s)
print(s_tokens)
```

```
['Please', 'pay', '$', '100.55', 'to', 'settle', 'your', 'bill.', 'Send', 'confirmatio  
n', 'to', 'confirm', '@', 'gmail.com', '.']
```

Part 4: Levenshtein Distance & Alignment

Relevant nltk documentation: <https://www.nltk.org/api/nltk.metrics.distance.html>

Exercise 4-1: Use the nltk functions `edit_distance` to compute the Levenshtein edit-distance between the strings "intention" and "execution"

```
In [35]: from nltk.metrics.distance import edit_distance  
  
w1 = "intention"  
w2 = "execution"  
dist = edit_distance(w1,w2)  
dist
```

```
Out[35]: 5
```

Exercise 4-2: Use the nltk function `edit_distance_align` to compute the minimum Levenshtein edit-distance based alignment mapping between the two strings "intention" and "execution"

```
In [38]: from nltk.metrics.distance import edit_distance_align  
  
w1 = "intention"  
w2 = "execution"  
dist_align = edit_distance_align(w1,w2)  
dist_align
```

```
Out[38]: [(0, 0),  
          (1, 1),  
          (2, 2),  
          (3, 3),  
          (4, 4),  
          (5, 5),  
          (6, 6),  
          (7, 7),  
          (8, 8),  
          (9, 9)]
```

```
In [37]: help(edit_distance_align)
```

Help on function `edit_distance_align` in module `nltk.metrics.distance`:

```
edit_distance_align(s1, s2, substitution_cost=1)  
    Calculate the minimum Levenshtein edit-distance based alignment  
    mapping between two strings. The alignment finds the mapping  
    from string s1 to s2 that minimizes the edit distance cost.  
    For example, mapping "rain" to "shine" would involve 2  
    substitutions, 2 matches and an insertion resulting in  
    the following mapping:  
    [(0, 0), (1, 1), (2, 2), (3, 3), (4, 4), (4, 5)]  
    NB: (0, 0) is the start state without any letters associated  
    See more: https://web.stanford.edu/class/cs124/lec/med.pdf
```

In case of multiple valid minimum-distance alignments, the
backtrace has the following operation precedence:

1. Substitute s1 and s2 characters

2. Skip s1 character
3. Skip s2 character

The backtrace is carried out in reverse string order.

This function does not support transposition.

```
:param s1, s2: The strings to be aligned
:type s1: str
:type s2: str
:type substitution_cost: int
:rtype: List[Tuple(int, int)]
```

In []: